Statistical Learning

Homework exercise 2

1. **Playing with quantile regression**
   Divide the competition data to 8000 training and 2000 validation observations, and fit linear quantile regression for quantiles $\{0.1, 0.25, 0.5, 0.75, 0.9\}$. Apply each model to the validation data, calculate and plot the following, and comment on the results:

   (a) Validation RMSE for each model.

   (b) Average prediction for each model.

   (c) The predictions of all models at five randomly selected validation observations.

   Tip: A function to run quantile regression (including absolute loss regression) in R is rq in package quantreg.

2. **Robustness of absolute loss regression and its variants compared to least squares**

   (a) **Theoretical:** Assume we have a sample $X = \{x_1, \ldots, x_n\}$. Given a statistic $S(X)$, the *breakdown point* is defined as the minimal % of samples that have to be contaminated (in any way we wish) so that $S(X)$ can be *arbitrarily* corrupted (i.e., driven to $\pm\infty$). Show that the breakdown point for the mean $S_1(X) = \bar{x}$ is $1/n$ whereas for the median $S_2(X) = \text{median}(X)$ the breakdown is about $1/2$. What would be the breakdown point for other quantiles?

   (b) **Practical:** Demonstrate the lack of robustness of linear regression compared to absolute loss regression by applying them both to the Prostate cancer dataset and plotting the effect that extreme contamination of one y value in the training set has on the test set RMSE, when each of these modeling approaches is used to calculate $\hat{\beta}$.

3. **Ways of interpreting and calculating Ridge and Lasso regression:**

   (a) **ESL 3.7:** Show that if we assume a likelihood $y_i \sim N(x_i^T\beta, \sigma^2)$ for $i = 1, ..., n$ and a prior $\beta \sim N(0, \tau^2 I)$, then the negative log-posterior density of β is $\sum_{i=1}^{n}(y_i - x_i^T\beta)^2 + \lambda \sum_{j=1}^{p} \beta_j^2$ up to multiplication and addition of constants, with $\lambda = \sigma^2/\tau^2$. Conclude that the Ridge solution is a maximum posterior estimate of β.

   (b) Show that the same applies to Lasso, except that the prior on β is a double exponential.
   Note: A double exponential random variable with parameter θ has density $f(x) = \theta/2 \cdot \exp(-|x|\theta)$.

   (c) **ESL 3.10 (3.12 in 2nd ed.):** Show that the ridge regression estimate can be obtained by ordinary least squares regression on an augmented data set. We augment the centered matrix X with p additional rows $\sqrt{\lambda}I_{p\times p}$, and augment y with p zeros. Comment briefly on how we can think of Ridge shrinkage as adding more "neutral" observations with 0 response.

   (d) What would be a corresponding case for the Lasso penalty, where the shrinkage can be accomplished by adding data and solving the same fitting problem?
   Hint: Think beyond squared error loss.

4. **The effect of identical predictors in Ridge and Lasso (3.28,3.29 in ESL 2nd ed.):** Assume we have a univariate model with one x variable and no intercept. We fit constrained ridge regression and lasso with a given constraint s on the norm (ℓ_2 norm squared for ridge, ℓ_1 norm for lasso). Now we add a second identical variable $x^* = x$ and refit the models with the same constraint. What happens to the coefficients $\hat{\beta}$ of both models? How does the two-dimensional solution to the new problem relate to the one-dimensional solution to the old one in each case? Is it unique? Assume the constraint is much smaller than the norm of the least squares solution, so it is tight.

   Hint: The behavior of ridge and lasso under this scenario is quite different. Since both predictors x, x^* are identical, a coefficient can be divided between them in different ways which give the same fit. Consider what different divisions do to the norm of the coefficient vector in each case, and use that to infer the optimal solution. You can also simulate to gain intuition.

5. **Short intuition questions:**

   (a) If I believe that only a small number of my variables are important, which one (or more) of these four regularization approaches should I use?

      i. Ridge
      ii. Lasso
      iii. Variable selection
      iv. PCA regression

   (b) Same question, except that now I believe that only a low-dimensional linear subspace of the span of my variables is important.